

References

Collaborated with Crystal Huang.

Question 1: k -th Smallest Element from Two Lists

Suppose you are given two *sorted* lists $A[1, \dots, n]$ and $B[1, \dots, m]$ of size n and m , respectively. Give an $O(\log k)$ algorithm to find the k -th smallest element in $A \cup B$, i.e., of the combination of the two arrays. For simplicity, you can assume $k \leq \min(m, n)$. Justify the correctness and running time of your algorithm.

Solution:

Algorithm. $\text{FINDINDEX}(A, B, k)$ where $A[1, \dots, n], B[1, \dots, m]$ are sorted lists and $k \in \mathbb{N}$ where $k \leq \min(n, m)$. That is, for all $1 \leq i < n$, we have $A[i] \leq A[i + 1]$, and for all $1 \leq j < m$, we have $B[j] \leq B[j + 1]$:

- if $n = 0$, then return $B[k]$;
- if $m = 0$, then return $A[k]$;
- if $k = 1$, then return $\min(A[1], B[1])$;
- otherwise, for $k > 1$, compute $i \leftarrow \lfloor \frac{k}{2} \rfloor$. Then:
 - if $A[i] \leq B[i]$, then return $\text{FINDINDEX}(A[i + 1, \dots, n], B[1, \dots, m], k - i)$;
 - otherwise, return $\text{FINDINDEX}(A[1, \dots, n], B[i + 1, \dots, m], k - i)$.

Proposition I.

Claim. For all sorted lists $A[1, \dots, n], B[1, \dots, m]$ and all $k \leq \min(n, m)$, $\text{FINDINDEX}(A, B, k)$ returns $C[k]$, where $C = \text{MERGE}(A, B)$, the result of the well-known MERGE algorithm. That is, $\text{FINDINDEX}(A, B, k)$ returns the k -th element of the sorted list that contains all elements of sorted lists A and B (and no other elements). Note $k \geq 1, m \geq 0, n \geq 0$.

Proof.

- Suppose $n = 0$. Then $m \geq 1$. Consider $C[1, \dots, m]$, the result of merging A and B . Since $n = 0$, we know $C = B$. For all $k \in \mathbb{N}$, $\text{FINDINDEX}(A, B, k)$ returns $B[k] = C[k]$. Thus, when $n = 0$, the claim holds for all $k \in \mathbb{N}$.
- Suppose $m = 0$. Then $n \geq 1$. Consider $C[1, \dots, n]$, the result of merging A and B . Since $m = 0$, we know $C = A$. For all $k \in \mathbb{N}$, $\text{FINDINDEX}(A, B, k)$ returns $A[k] = C[k]$. Thus, when $m = 0$, the claim holds for all $k \in \mathbb{N}$.
- Suppose instead $m > 0$ and $n > 0$. We can demonstrate the claim for $m > 0$ and $n > 0$ by induction on k .

Basis. Consider $k = 1$. Consider also $C[1, \dots, m + n]$, the result of merging A and B . The least element in C is either the least element in A or the least element in B ; specifically, it

is the smaller of the least element in A and the least element in B . Of course, in this case $\text{FINDINDEX}(A, B, k)$ returns $\min(A[1], B[1])$. Since A, B are sorted lists, $A[1]$ is the least element in A and $B[1]$ is the least element in B . Therefore $\text{FINDINDEX}(A, B, k)$ returns the first smallest element in C . Since $k = 1$, the basis holds for all $A[1, \dots, n], B[1, \dots, m]$.

Hypothesis. Consider $1 < k \leq \ell < \min(n, m)$. Assume that for all sorted lists $A[1, \dots, n]$ and $B[1, \dots, m]$, $\text{FINDINDEX}(A, B, \ell)$ returns the ℓ -th element of the sorted list that contains all elements of A and B (and no other elements).

Inductive step. Consider $k = \ell + 1$. For all sorted lists $A[1, \dots, n], B[1, \dots, m]$, we have $i = \lfloor \frac{\ell+1}{2} \rfloor$. Consider $C[1, \dots, m+n]$, the result of merging A and B .

- Suppose $A[i] \leq B[i]$. Since A and B are sorted, and since $A[i] \leq B[i]$, we know that the $(\ell + 1)$ -th element in C is definitely not within the first i elements of A . We can ignore the first i elements in A , adjusting our index in the subarray from $\ell + 1$ to $\ell + 1 - i$ (to account for the i elements excluded, which precede it in sorted order).

The result for the new subarray is given by $\text{FINDINDEX}(A[i+1, \dots, n], B[1, \dots, m], k-i)$, which is the return value of the algorithm in this case. By the strong induction hypothesis, the claim holds for this result.

- Suppose instead $A[i] > B[i]$. Since A and B are sorted, and since $B[i] < A[i]$, we know that the $(\ell + 1)$ -th element in C is definitely not within the first i elements of B . We can ignore the first i elements in B , adjusting our index in the subarray from $\ell + 1$ to $\ell + 1 - i$ (to account for the i elements excluded, which precede it in sorted order).

The result for the new subarray is given by $\text{FINDINDEX}(A[1, \dots, n], B[i+1, \dots, m], k-i)$, which is the return value of the algorithm in this case. By the strong induction hypothesis, the claim holds for this result.

In all cases, $\text{FINDINDEX}(A, B, \ell + 1)$ returns the $(\ell + 1)$ -th element of the sorted list that contains all elements of A and B (and no other elements), thus completing the inductive step.

Hence, by the principle of mathematical induction, the claim holds for all $k \in \mathbb{N}$ when $m > 0$ and $n > 0$.

We have shown that for all sorted lists $A[1, \dots, n], B[1, \dots, m], n \geq 0, m \geq 0, k \in \mathbb{N}$, if $k \leq \min(n, m)$, then $\text{FINDINDEX}(A, B, k)$ returns the k -th element of the sorted list that contains all elements of sorted lists A and B . \square

Proposition II.

Claim. For all sorted lists $A[1, \dots, n], B[1, \dots, m]$ and all $k \leq \min(n, m)$, the running time of $\text{FINDINDEX}(A, B, k)$ is $T(k) = O(\log k)$.

Proof. On each recursive level, $\text{FINDINDEX}(A, B, k)$ performs a constant number of operations: At most, one element from A is compared with one element from B . So the work done on each level is $\Theta(1)$. Since k is reduced by roughly half on each recursive call, the number of levels is asymptotically $\log_2 k$. From this intuition, we can guess that the running time of $\text{FINDINDEX}(k)$, $T(k)$ may be $O(\log_2 k)$.

We can use a recurrence to estimate $T(k)$:

$$T(1) = \Theta(1), \quad T(k) = T\left(\left\lfloor \frac{k}{2} \right\rfloor\right) + \Theta(1), \text{ for } k > 1.$$

We want to show that $T(k) = O(\log_2 k)$. That is, we want to show that there exist positive constants C, n_0 such that $T(k) \leq C \log k$ for all $k \geq n_0$. Taking $k_0 = 2$, we can demonstrate this claim for all $k > 2$ by induction on k .

Basis. Consider $k = 2$. Observe

$$\begin{aligned} T(2) &= T\left(\left\lfloor \frac{2}{2} \right\rfloor\right) + \Theta(1) \\ &= T(1) + \Theta(1) \\ &= \Theta(1) + \Theta(1) \\ &= c_1 + c_2 \end{aligned} \quad \begin{array}{l} \text{asymptotically,} \\ \text{where } c_1, c_2 \text{ are positive constants.} \end{array}$$

Since $T(2) = c_1 + c_2$, we can choose any $C \geq \frac{c_1+c_2}{\log_2 k}$ so that $T(2) \leq C \log_2 k$. The claim holds in the base case.

Hypothesis. Consider $2 < \ell \leq k$. Assume that $T(\ell) = O(\log_2 \ell)$.

Inductive step. Consider $k = \ell + 1$. Observe

$$\begin{aligned} T(\ell + 1) &= T\left(\left\lfloor \frac{\ell + 1}{2} \right\rfloor\right) + \Theta(1) \\ &= O\left(\log_2 \left\lfloor \frac{\ell + 1}{2} \right\rfloor\right) + \Theta(1) && \text{from the strong inductive hypothesis,} \\ &\leq C \log_2 \left(\left\lfloor \frac{\ell + 1}{2} \right\rfloor\right) + \Theta(1) \\ &\leq C \log_2 \left(\left\lfloor \frac{\ell + 1}{2} \right\rfloor\right) + \Theta(1) \\ &\leq C \log_2 \ell + \Theta(1) && \text{for } C > 0 \text{ and } \ell \geq 2, \\ &= O(\log_2 \ell) + \Theta(1) \\ &= O(\log_2 \ell) && \text{asymptotically.} \end{aligned}$$

This completes the inductive step.

Hence, by the principle of mathematical induction, we have $k_0 = 2, C \geq \frac{c_1+c_2}{\log_2 k}$ for which $T(k) \leq C \log_2 k$ for all $k \geq k_0$.

Therefore, $T(k) = O(\log_2 k) = O(\log k)$. \square

Question 2: Permutations

Define the notation $[n] = \{1, 2, \dots, n\}$, and let S_n be the set of all possible permutations of $[n]$. The size of S_n is given by $|S_n| = n! = n \cdot (n-1) \cdot \dots \cdot 1$. Recall that $n! = O(n^n)$ and $2^n = O(n!)$. Now, each input in S_n can serve as an input for a sorting algorithm. Instead of a perfectly correct sorting algorithm, we will look at a class of algorithms that only produce a sorted result on some of the inputs. More concretely, we say that a sorting algorithm is ε -correct if the algorithm produces the correct result (i.e., produces a sorted array as output) on exactly ε fraction of the set of inputs in S_n . In other words, an ε -correct sorting algorithm is one that produces the correct result on $\varepsilon \cdot (n!)$ possible inputs and produces an incorrect result otherwise.

1. Show that for any $0 \leq \varepsilon \leq 1$, the decision tree of an ε -correct comparison-based sorting algorithm must have at least $\varepsilon \cdot n!$ leaves.

Solution: *Claim.* Let $0 \leq \varepsilon \leq 1$. Then the decision tree of an ε -correct comparison-based sorting algorithm must have at least $\varepsilon \cdot n!$ leaves.

Proof. Assume, for the sake of contradiction, that there exists some ε -correct sorting algorithm `HYPOTHETICALSORT` for which the number of leaves is less than $\varepsilon \cdot n!$. Since `HYPOTHETICALSORT` is an ε -correct sorting algorithm, it produces correct solutions for $\varepsilon \cdot n!$ possible inputs. Every input array A has exactly one correct sorted solution A_{sorted} . Since `HYPOTHETICALSORT` produces correct results for $\varepsilon \cdot n!$ possible inputs, it produces exactly $1 \cdot \varepsilon \cdot n! = \varepsilon \cdot n!$ possible outputs. In a recursive decision tree, each leaf corresponds to a single output. Therefore, the decision tree of `HYPOTHETICALSORT` has $\varepsilon \cdot n!$ leaves. This contradicts the hypothesis that `HYPOTHETICALSORT` has less than $\varepsilon \cdot n!$ leaves.

Therefore for all $0 \leq \varepsilon \leq 1$, the decision tree of an ε -correct comparison-based sorting algorithm must have at least $\varepsilon \cdot n!$ leaves. \square

In the following, we want to investigate whether lowering ε can yield a saving in the number of comparisons required by a sorting algorithm. Intuitively if the algorithm only has to be correct on a certain fraction of inputs this should speed up the algorithm. For instance, an algorithm that does not need to be correct at all clearly does not need $\Omega(n \log n)$ comparisons. (**Hint:** Use the number of leaves in the decision tree to derive a lower bound on its height.)

2. Let $\varepsilon = 1/2$. Show that for any constant $C \geq 0$, there is no comparison-based ε -correct sorting algorithm that can sort using less than $C \cdot n$ comparisons. This shows that taking $\varepsilon = 1/2$ does not help us reduce the number of comparisons to linear.

Solution: *Claim.* Let $\varepsilon = \frac{1}{2}$. Then for any $C \geq 0$, there is no comparison-based ε -correct sorting algorithm that can sort using less than Cn comparisons.

Proof. Assume, for the sake of contradiction, that there exists some comparison-based ε -correct sorting algorithm `HYPOTHETICALSORT` that can sort using less than Cn comparisons for $\varepsilon = \frac{1}{2}$. Then the number of leaves in its decision tree is at least $\varepsilon \cdot n! = \frac{1}{2} \cdot n! = \frac{n!}{2}$. Thus, the height of this tree $H(n)$ is at least on the order of $\log \frac{n!}{2}$. Note

$$\frac{n!}{2} \geq \frac{\left(\frac{n}{2}\right)^{\frac{n}{2}}}{2}.$$

Observe

$$\begin{aligned} H(n) &\geq \log \frac{n!}{2} \\ &\geq \log \frac{\left(\frac{n}{2}\right)^{\frac{n}{2}}}{2} \\ &\geq \log \left(\frac{n}{2}\right)^{\frac{n}{2}} - \log 2 \\ &\geq \frac{n}{2} \log \frac{n}{2} - \log 2. \end{aligned}$$

We have $C_1 \left(\frac{n}{2} \log \frac{n}{2} - \log 2\right) \geq n \log n$ for all $n \geq n_0$, taking $C_1 > 3$ and large enough n_0 .

Thus, $H(n) = \Omega(n \log n) = \Omega(n)$. Therefore `HYPOTHETICALSORT` sorts using greater than or equal to Cn comparisons for some $C = C_1$. This contradicts the hypothesis that `HYPOTHETICALSORT` can sort using less than Cn comparisons for all $C \geq 0$.

Therefore, when $\varepsilon = \frac{1}{2}$, for all $C \geq 0$, there is no comparison-based ε -correct sorting algorithm that can sort using less than Cn comparisons. \square

3. Consider $\varepsilon = 1/n$. In this setting, are we able to achieve a sorting algorithm for S_n with $O(n)$ comparisons? Justify your answer.

Solution: *Claim.* Let $\varepsilon = \frac{1}{n}$. Then for any $C \geq 0$, there is no comparison-based ε -correct sorting algorithm that can sort using $O(n)$ comparisons—that is, less than or equal to Cn comparisons.

Proof. Assume, for the sake of contradiction, that there exists some comparison-based ε -correct sorting algorithm `HYPOTHETICALSORT` that can sort using less than or equal to Cn comparisons for $\varepsilon = \frac{1}{n}$. Then the number of leaves in its decision tree is at least $\varepsilon \cdot n! = \frac{1}{n} \cdot n! = (n-1)!$. Thus, the height of this tree $H(n)$ is at least on the order of $\log((n-1)!)$. Note

$$(n-1)! \geq \left(\frac{n-1}{2}\right)^{\frac{n-1}{2}}.$$

Observe

$$\begin{aligned} H(n) &\geq \log((n-1)!) \\ &\geq \log\left(\frac{n-1}{2}\right)^{\frac{n-1}{2}} \\ &\geq \frac{n-1}{2} \log \frac{n-1}{2}. \end{aligned}$$

We have $C_1 \left(\frac{n-1}{2} \log \frac{n-1}{2}\right) > n \log n$ for all $n \geq n_0$, taking $C_1 \geq 3$ and large enough n_0 .

Thus, $H(n) = \omega(n \log n)$. We know that $\omega(n \log n) = \omega(n)$ is strictly greater than $O(n)$ for large enough n . Therefore `HYPOTHETICALSORT` sorts using greater than Cn comparisons for some $C = C_1$. This contradicts the hypothesis that `HYPOTHETICALSORT` can sort using $O(n)$ comparisons.

Therefore, when $\varepsilon = \frac{1}{n}$, for all $C \geq 0$, there is no comparison-based ε -correct sorting algorithm that can sort using $O(n)$ comparisons. \square

4. Consider $\varepsilon = \frac{1}{2^n}$. In this setting, are we able to achieve a sorting algorithm for S_n with $O(n)$ comparisons? Justify your answer.

Solution: *Claim.* Let $\varepsilon = \frac{1}{2^n}$. Then for any $C \geq 0$, there is no comparison-based ε -correct sorting algorithm that can sort using $O(n)$ comparisons—that is, less than or equal to Cn comparisons.

Proof. Assume, for the sake of contradiction, that there exists some comparison-based ε -correct sorting algorithm `HYPOTHETICALSORT` that can sort using less than or equal to Cn comparisons for $\varepsilon = \frac{1}{2^n}$. Then the number of leaves in its decision tree is at least $\varepsilon \cdot n! = \frac{1}{2^n} \cdot n! = \frac{n!}{2^n}$. Thus, the height of this tree $H(n)$ is at least on the order of $\log\left(\frac{n!}{2^n}\right)$. Note

$$\frac{n!}{2^n} \geq \frac{\left(\frac{n}{2}\right)^{\frac{n}{2}}}{2^n}.$$

Observe

$$\begin{aligned} H(n) &\geq \log\left(\frac{n!}{2^n}\right) \\ &\geq \log\frac{\left(\frac{n}{2}\right)^{\frac{n}{2}}}{2^n} \\ &\geq \log\left(\frac{n}{2}\right)^{\frac{n}{2}} - \log 2^n \\ &= \frac{n}{2} \log \frac{n}{2} - n \log 2. \end{aligned}$$

We have $C_1 \left(\frac{n}{2} \log \frac{n}{2} - n \log 2\right) > n \log n$ for all $n \geq n_0$, taking $C_1 \geq 10$ and large enough n_0 . Thus, $H(n) = \omega(n \log n)$. We know that $\omega(n \log n) = \omega(n)$ is strictly greater than $O(n)$ for large enough n . Therefore `HYPOTHETICALSORT` sorts using greater than Cn comparisons for some $C = C_1$. This contradicts the hypothesis that `HYPOTHETICALSORT` can sort using $O(n)$ comparisons.

Therefore, when $\varepsilon = \frac{1}{2^n}$, for all $C \geq 0$, there is no comparison-based ε -correct sorting algorithm that can sort using $O(n)$ comparisons. \square

5. Consider $\varepsilon = \frac{2^n}{n!}$. In this setting, are we able to achieve a sorting algorithm for S_n with $O(n)$ comparisons?

Solution: *Claim.* Let $\varepsilon = \frac{2^n}{n!}$. Then for any $C \geq 0$, there exists a comparison-based ε -correct sorting algorithm that can sort using $O(n)$ comparisons—that is, less than or equal to Cn comparisons.

Proof. Let `Sort` be a comparison-based ε -correct sorting algorithm that can sort using less than or equal to Cn comparisons for $\varepsilon = \frac{2^n}{n!}$. Then the number of leaves in its decision tree is at least $\varepsilon \cdot n! = \frac{2^n}{n!} \cdot n! = 2^n$. Thus, the height of this tree $H(n)$ is at least on the order of $\log 2^n = n \log 2$. Observe

$$H(n) \geq n \log 2.$$

We have $n \log 2 \leq Cn$ for $C \geq \log 2$ for all $n \geq 0$. Thus, $H(n) = O(n)$, so `Sort` can sort using $O(n)$ comparisons. \square

Question 3: Disjointed arrays

Let $A[1, \dots, m]$ and $B[1, \dots, n]$ be two sorted arrays each containing distinct elements. Let $m \leq n$ and n is a multiple of m . The problem is to determine if the two arrays are disjoint or not. Two arrays are said to be disjoint if their intersection is \emptyset .

1. Let us assume that A and B are disjoint. We define C as the sorted combination of A and B , i.e., $C = A \cup B$ and C is sorted. Clearly, $|C| = m + n$ because $A \cap B = \emptyset$. Define D of length $m + n$ where $D[i] = 1$ if $C[i] \in A$ and 0 otherwise. Give a count of the number of such possible arrays D . Justify your answer.

Solution: Proposition 3.1. *Claim.* Let $A[1, \dots, m], B[1, \dots, n]$ be sorted arrays containing distinct elements where $m \leq n$ and n is a multiple of m . Let $C = A \cup B$ where C is sorted and $D[1, \dots, m + n]$ where $D[i] = 1$ if $C[i] \in A$, and $D[i] = 0$ otherwise, for $1 \leq i \leq m + n$.

Proof. Since $C = A \cup B$, we know $D[i] = 1$ if $C[i] \in A$ or 0 if $C[i] \in B$. There are $m + n$ many elements in D , and for each element, it is either in A or not in A . It is known that there are n elements in B , so n of $m + n$ elements must be chosen to be 0 and the remaining m elements must be 1.

Thus, the number of such possible arrays is $\binom{m+n}{n} = \binom{m+n}{m}$. \square

2. We go back to the original problem—we do not know whether A, B are disjoint. Let us assume that A contains 1 element, and B contains 2 elements.
 - Draw a comparison-based decision tree for this problem.
 - You will also present a modification of the above tree where you will change the label of every leaf node marked as true with a corresponding array D .

You will ensure that your decision tree makes the least number of comparisons possible and has the shortest height possible. The internal node will be of the form (i, j) which indicates that you are comparing $A[i]$ and $B[j]$. Now, each such internal node will have three children—one corresponding to $A[i] < B[j]$, one corresponding to $A[i] = B[j]$, and one corresponding to $A[i] > B[j]$. The leaf nodes will contain the values true, false where true indicates that A, B are disjoint and false indicates the opposite.




Figure 1: *Top*: A comparison-based decision tree for this problem. *Bottom*: A modification of the above tree where the label of every leaf node marked as true is replaced with a corresponding array D .

Solution: Please see Figure 1.

3. Now we look at the general problem, i.e., for any m, n . It can be shown that the decision tree for this general problem will have every leaf node labeled true correspond with a distinct array D , as defined in part (1). This is in fact a bijective mapping, i.e., every leaf node labeled true has a corresponding array D and every array D has a corresponding leaf node labeled true. Now, using this and your answer from part (1), show that problem has a worst-case lower bound of $\Omega(m \log(1 + n/m))$.

Solution: Proposition 3.3. Definitions. Let $A[1, \dots, n], B[1, \dots, m]$ be sorted arrays containing distinct elements where $m \leq n$ and n is a multiple of m . Let $C = A \cup B$ where C is sorted. Let $D[1, \dots, m+n]$ be an array where $D[i] = 1$ if $C[i] \in A$ and 0 otherwise.

Axiom I. There exists a bijective mapping between nodes labeled true and possible arrangements of D . For each node labeled true, there is exactly one arrangement of D . For each arrangement of D , there is exactly one node labeled true.

Axiom II. $\binom{a}{b} \geq \left(\frac{a}{b}\right)^b$ for $a \geq b \geq 1$. This follows from $b^b \geq b!$ for all $b \geq 1$.

Claim. The problem has a worst-case lower bound of $\Omega\left(m \log\left(1 + \frac{n}{m}\right)\right)$.

Proof. Let $H(n, m)$ denote the height of the decision tree in the worst case. Since each internal node has three children—one corresponding to $A[i] < B[j]$, one corresponding to $A[i] = B[j]$, and one corresponding to $A[i] > B[j]$ —this is a ternary tree with height equal to the ternary logarithm of the number of leaves.

From Proposition 3.1, we know there are $\binom{m+n}{m}$ arrangements of D . From Axiom I, there are as many arrangements of D as nodes labeled true. Therefore there are $\binom{m+n}{m}$ nodes labeled true.

The worst case occurs when A and B are disjoint. Suppose not: Then $A[i] = B[j]$ for some i, j , and the result is false. We can choose different inputs A, B' , where $B'[1, \dots, m] = [B[1], \dots, B[j-1], x, B[j+1], \dots, B[m]]$ and $x \neq A[i]$. In other words, we can replace $B'[j]$ with some x to make A and B' disjoint and perform the algorithm again on A and B' . Now $A[i] \neq B'[j]$, so at least one additional comparison is performed. The running time will be as bad, or worse, than when A and B are not disjoint. So the worst case is when A and B are disjoint.

When A and B are disjoint, the result is true. There are $\binom{m+n}{m}$ nodes labeled true, so in the worst case all $\log_3 \binom{m+n}{m}$ levels must be visited before the last evaluation concludes that A and B are disjoint. Intuitively, this corresponds to the case where the first $m-1$ elements compared between A and B are different, but the last element compared is identical.

Observe:

$$\begin{aligned} H(n, m) &\geq \log_3 \binom{m+n}{m} \\ &\geq \log_3 \left(\frac{m+n}{m} \right)^m && \text{from Axiom II, since } n \geq m \geq 1, \\ &\geq \log_3 \left(1 + \frac{n}{m} \right)^m \\ &\geq m \log_3 \left(1 + \frac{n}{m} \right) \\ &\geq \frac{1}{\log 3} \cdot m \log \left(1 + \frac{n}{m} \right). \end{aligned}$$

There exists a constant $\frac{1}{\log 3} > 0$ for which $H(n, m) \geq \frac{1}{\log 3} \cdot m \log \left(1 + \frac{n}{m} \right)$ for large enough m (and, since $n \geq m$, large enough n).

Therefore, $H(n, m) = \Omega(m \log(1 + n/m))$. Of course, the work done on each level is constant.

Ergo the problem has a worst-case lower bound of $\Omega(m \log(1 + n/m))$. \square