

## References

Peer-reviewed with Crystal Huang.

## Question 1: Breadth-first search

Recall the BFS algorithm from the lecture that runs on graph  $G$  starting from root node  $s \in V$ :

```
BFS( $G, s$ )
1    $s.color = gray, s.dist = 0.$ 
2   for  $u \in V \setminus \{s\}$ 
3        $u.color = white, u.dist = \infty$ 
4    $Q = \emptyset$ 
5    $Q.enqueue(s)$ 
6   while  $Q \neq \emptyset$ 
7        $u = Q.dequeue()$ 
8       for  $v \in Adj[u]$ 
9           if  $v.color = white$ 
10                $v.color = gray$ 
11                $v.dist = u.dist + 1$ 
12                $Q.enqueue(v)$ 
13    $u.color = black$ 
```

1. Consider the undirected graph shown in Figure 1. If we begin a BFS traversal starting at node  $A$ , in what order are the nodes visited? Assume that the adjacency list is ordered alphabetically, i.e., we visit earlier letters alphabetically first, so from  $A$  we will visit  $B$  before  $C$ , and so on.




Figure 1: Graph to traverse

**Solution:** We can step through the breadth-first search algorithm to determine that the order of the vertices visited is  $A, B, C, E, F, K, D, H, G, I, L, M$ , and  $J$ .

*First.* We enqueue  $A$ .

*Second.* We dequeue  $A$  and enqueue  $B$  and  $C$ .

*Third.* We dequeue  $B$  and enqueue  $E$ .

*Fourth.* We dequeue  $C$  and enqueue  $F$  and  $K$ .

*Fifth.* We dequeue  $E$  and enqueue  $D$  and  $H$ .

*Sixth.* We dequeue  $F$ . We skip  $H$ . We enqueue  $G$  and  $I$ .

*Seventh.* We dequeue  $K$ .

*Eighth.* We dequeue  $H$  and enqueue  $L$  and  $M$ .

*Ninth.* We dequeue  $G$  and enqueue  $J$ .

*Tenth.* We dequeue  $I$ . We skip  $J$ .

*Eleventh, twelfth, and thirteenth.* We dequeue  $L, M$ , and  $J$ .

The vertices are visited as enqueued:  $A, B, C, E, F, K, D, H, G, I, L, M$ , and  $J$ .

2. Explain what happens if we don't check whether a node is white before enqueueing it.

**Solution:** Suppose we did not check whether a node is white before enqueueing. Then the breadth-first search algorithm may not terminate.

For example, suppose we perform this variant of breadth-first search on graph  $G = (V, E)$  with vertices  $V = \{A, B, C\}$  and edges  $E = \{\{A, B\}\}$ . Assume, without loss of generality, that the initial vertex is A and that other vertices are visited in alphabetical order.

We can step through the algorithm to demonstrate the flaw.

*First.* We enqueue A.

*Second.* We dequeue A and enqueue B.

*Third.* We dequeue B and enqueue A. Note that we enqueued A in this situation because we did not check its color and thus ignored the fact that it had been previously visited.

This brings us back to the second step (dequeue A and enqueue B), meaning that the algorithm never terminates, even after all vertices have been visited.

3. Assume we replace the (FIFO) queue with a stack and therefore in line 7 obtain the node *last* added. Does  $v.dist$  still represent the shortest distance from  $s$  to  $v$ ?

**Solution:** No, using this implementation of breadth-first search, the distance assigned to a vertex  $v$  may not represent the shortest distance from the starting vertex  $s$  to  $v$ .

Consider the undirected graph  $G = (V, E)$  with vertices  $V = \{A, B, C, D, E\}$  and edges  $E = \{\{A, B\}, \{A, C\}, \{B, E\}, \{C, D\}, \{D, E\}\}$ .

Assume, without loss of generality, that the initial vertex is A and that other vertices are visited in alphabetical order.

We can see that the shortest path from A to E is  $(A, B, E)$ , so the shortest distance from A to E is 2.

However, if we step through the modified breadth-first search algorithm (a hybrid of breadth-first search and depth first search), we see that the algorithm handles the immediate children of the source node correctly, but mishandles its other descendent vertices.

*First.* We push A with distance 0.

*Second.* We pop A, noting edges  $\{A, B\}$  and  $\{A, C\}$ . We push B with distance  $0 + 1 = 1$  and C with distance  $0 + 1 = 1$ .

*Third.* We pop C, noting edge  $\{C, D\}$ . We push D with distance  $1 + 1 = 2$ .

*Fourth.* We pop D, noting edge  $\{D, E\}$ . We push E with distance  $2 + 1 = 3$ . Already, this result is incorrect. We can continue stepping through the algorithm to show that it remains incorrect at termination.

*Fifth.* We pop E.

*Sixth.* We pop B, noting edge  $\{B, E\}$ . We skip E.

The algorithm incorrectly assigns a distance of 3 to E, although the shortest distance from A to E is 2. We conclude that, in the modified algorithm, the distance assigned to  $v$  is not necessarily the shortest distance from  $s$  to  $v$ .

4. Consider graphs for which edges have assigned weights. Assume we modify line 11 of BFS to add the weight of the edge from  $u$  to  $v$  instead of adding 1. Does the modified algorithm compute the shortest weighted distance correctly? If so, prove the correctness; otherwise, provide a counterexample for which the algorithm fails.

**Solution:** No, the modified algorithm is incorrect.

For example, suppose we perform this variant of the algorithm on graph  $G = (V, E)$  with vertices  $V = \{A, B, C, D\}$  and edges  $E$ . For each edge  $e \in E$ , note  $e = (u, v, w)$ , where  $u$  is the source vertex,  $v$  is the target vertex, and  $w$  is the weight of the edge. Suppose

$$E = \{(A, B, 10), (A, C, 1), (B, D, 20), (C, D, 2)\}.$$

Assume, without loss of generality, that  $G$  is directed, that the initial vertex is A, and that other vertices are visited in alphabetical order.

We can step through the algorithm to demonstrate the flaw.

*First.* We enqueue A with distance 0.

*Second.* We dequeue A, noting edges  $(A, B, 10)$  and  $(A, C, 20)$ . We enqueue B with distance  $0 + 10 = 10$  and C with distance  $0 + 1 = 1$ .

*Third.* We dequeue B, noting edge  $(B, D, 20)$ . We enqueue D with distance  $10 + 20 = 30$ .

*Fourth.* We dequeue C, noting edge  $(C, D, 2)$ . We skip D.

This gives 30 for the distance from A to D by following path  $(A, B, D)$ . However, there exists a path  $(A, C, D)$  whereby the distance from A to D is 3. Of course,  $3 < 30$ . Thus, the modified algorithm fails.

## Question 2: Universal sink

Recall that the *adjacency matrix* representation of a *directed* graph  $G = (V, E)$  is the  $|V| \times |V|$  matrix  $M_G$  where  $M_G[i][j]$  is 1 if there is an edge from vertex  $i$  to vertex  $j$ , and 0 otherwise. You may assume no vertex in the graph has an edge going into itself (i.e., no self-loops). We want to determine whether there is any vertex in  $G$  that has edges coming to it from *all other* vertices but no edges going out from it (this is known as a **universal sink**).

1. Write an  $O(|V|^2)$  algorithm that, given  $M_G$ , checks if there is a universal sink.

### Solution:

**Algorithm I.**  $\text{IsSink}(M_G, v)$  with a  $|V| \times |V|$  adjacency matrix  $M_G$  representing directed graph  $G = (V, E)$  with no self-loops and a vertex  $v \in V$ ; returns true if  $v$  is a universal sink; otherwise, false:

For  $u \in (V \setminus \{v\})$ , if  $M_G[v, u] = 1$  or  $M_G[u, v] = 0$ , then return false.

Return true.

**Algorithm II.**  $\text{QuadraticSink}(M_G)$  with a  $|V| \times |V|$  adjacency matrix  $M_G$  representing directed graph  $G = (V, E)$  with no self-loops; returns true if there is a universal sink in  $G$ ; otherwise, false:

For  $v \in V$ , if  $\text{IsSink}(M_G, v)$  is true, then return true.

Return false.

**Proposition I.** *Claim.* Algorithm I determines if  $v$  is a universal sink in  $G$  in running time  $O(|V|)$ .

*Proof.* By definition,  $v$  is a universal sink in  $G = (V, E)$  if, for all  $u \in V$  where  $u \neq v$ , we have  $(u, v) \in E$  but  $(v, u) \notin E$ . The  $\text{IsSink}$  algorithm compares  $v$  to all other vertices  $u$ . It returns true only if  $v$  has edges coming to it from all other vertices  $u$  and no edge going out from it to any vertex  $u$ . Otherwise,  $\text{IsSink}$  returns false. Thus, Algorithm I correctly determines if  $v$  is a universal sink.

This implementation of Algorithm I visits every element in  $V$  and performs constant-time operations within the loop body, so it has running time  $O(|V| \times 1) = O(|V|)$ .  $\square$

**Proposition II.** *Claim.* Algorithm II determines if there exists a universal sink in  $G$  in running time  $O(|V|^2)$ .

*Proof.* For  $G = (V, E)$ , The  $\text{QuadraticSink}$  algorithm checks every vertex  $v \in V$ , using Algorithm I to determine if  $v$  is a universal sink. From Proposition I, Algorithm I correctly determines if  $v$  is a universal sink. Since Algorithm II returns true if Algorithm I determines that any vertex is a universal sink (and false if none satisfies this test), Algorithm II correctly determines if there exists a universal sink in  $G$ .

This implementation of Algorithm II visits every element in  $V$  and performs  $\text{IsSink}$  on each ( $O(|V|)$  by Proposition I) within the loop body, so it has running time  $O(|V| \times |V|) = O(|V|^2)$ .  $\square$

2. Write an  $O(|V|)$  algorithm for the problem. Justify why it is correct, as well as why it satisfies this run-time bound.

**Solution:**

**Algorithm III.** `LINEAR_SINK( $M_G$ )` with a  $|V| \times |V|$  adjacency matrix  $M_G$  representing directed graph  $G = (V, E)$  with no self-loops; returns `true` if there is a universal sink in  $G$ ; otherwise, `false`:

If  $V = \emptyset$ , then return `false`.

Let  $v^* \leftarrow v \in V$ , choosing  $v$  arbitrarily.

Let  $U \leftarrow V \setminus \{v^*\}$ .

Let  $W \leftarrow \emptyset$ .

For  $u \in U$ :

- if  $M_G[v^*][u] = 1$  then:
  - assign  $W \leftarrow W \cup \{v^*\}$ ;
  - assign  $v^* \leftarrow u$ ;
- otherwise:
  - assign  $W \leftarrow W \cup \{u\}$ .

Return `IsSink( $M_G, v^*$ )`.

**Invariant I. Claim.** For every step of the loop, no vertex  $w \in W$  is a universal sink.

*Proof.* We can prove this invariant by structural induction on  $W$ .

*Basis.* Consider  $W = \emptyset$ . Then the invariant holds vacuously at initialization.

*Hypothesis.* Consider  $W = X$  where  $X \neq \emptyset$ . Assume there exists no universal sink  $w \in W$ .

*Inductive step.* Consider  $W = X \cup \{v\}$  for  $v \in V$ . From Algorithm III, there are two cases:

- Suppose  $M_G[v^*][u] = 1$ . Then  $v = v^*$  (before the reassignment of  $v^*$ ). Thus  $M_G[v][u] = 1$ , indicating that there is an outgoing edge from  $v$ , or  $(v, u) \in E$ . So  $v$  is not a universal sink.
- Suppose instead  $M_G[v^*][u] = 0$ . Then  $v = u$ . Thus  $M_G[v^*][v] = 0$ , indicating that there is no incoming edge to  $v$ , or  $(v^*, v) \notin E$ . So  $v$  is not a universal sink.

In all cases,  $v$  is not a universal sink.

For all  $x \in X$ , the invariant holds by the inductive hypothesis. Since the invariant holds for  $v$ , it must therefore hold for all vertices  $w \in (X \cup \{v\})$ .

Hence, by structural induction, there exists no universal sink  $w \in W$ . This invariant holds at initialization, maintenance, and termination.

**Proposition III.** *Claim.* Algorithm III determines if there exists a universal sink in  $G$  in running time  $O(|V|)$ .

*Proof.* Suppose  $V = \emptyset$ . Then there is no universal sink in  $G$ , and Algorithm III returns false, which is correct.

Suppose instead  $V \neq \emptyset$ . Then there exists an arbitrary  $v^* \in V$ .

The loop executes for each of the  $(|V| - 1)$  elements in  $U = V \setminus \{v^*\}$ .

From Algorithm III, at initialization we have  $|W| = 0$ . Also, for each iteration, exactly one element is inserted into  $W$  in all cases. This implies that  $|W| = (|V| - 1)$  at termination. From Invariant I, at termination, we have for all  $w \in W$  that  $w$  is not a universal sink. Of course  $v^* \notin W$ , but  $v^* \in V$ . In fact,  $v^*$  is the only vertex that has not been eliminated.

If  $v^*$  is a universal sink, then there is a universal sink in  $G$ ; otherwise, there is none.

Algorithm III returns  $\text{IsSink}(M_G, v^*)$ , which, by Proposition I, correctly determines if  $v^*$  is a universal sink. Thus, Algorithm III correctly determines if there exists a universal sink in  $G$ .

The loop body performs constant-time operations in all cases, so the running time of the loop is  $O((|V| - 1) \times 1) = O(|V| - 1)$ . The running time of  $\text{IsSink}(M_G, v^*)$  is  $O(|V|)$  from Proposition I. Thus, the running time of  $\text{LINEAR\_SINK}$  is  $O(|V| - 1) + O(|V|) = O(|V|)$ .  $\square$